

COMP3311 24T1

Database Systems

week 2 - 1



Outline



- **Announcements**
- ER-Relational Mapping
- SQL Introduction



Announcement 1



- Quiz 1
 - [Activities | COMP3311 24T1 | WebCMS3 \(unsw.edu.au\)](#)
 - You can try as many times as possible
 - before midnight Friday, 23rd Feb (11:59:59 pm)



Announcement 2



- Exam Hurdle: You need to score 40% in the final exam & have a 50% overall score to pass this course.
- Lecturing Hurdle: some other courses using this LT have the same issue.
 - COMP1521, MATH1081



Outline



- Announcements
- **ER-Relational Mapping**
- SQL Introduction



ER vs Relational Models – Recap

Correspondence between ER and Relational models:


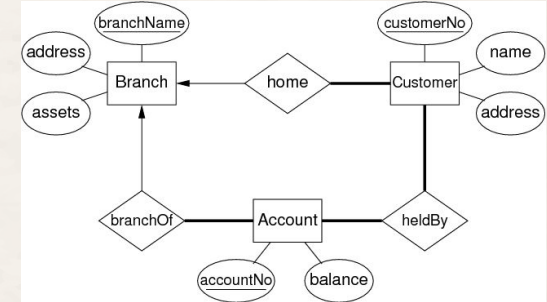
Relational attributes correspond to **ER attributes**

- although ER attributes generally don't have explicit domains

Relational tuples correspond to **ER entities**

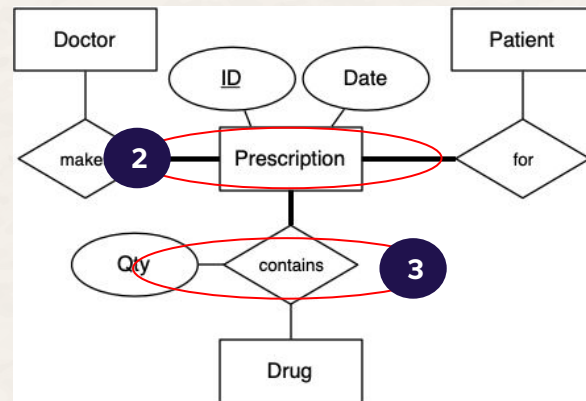
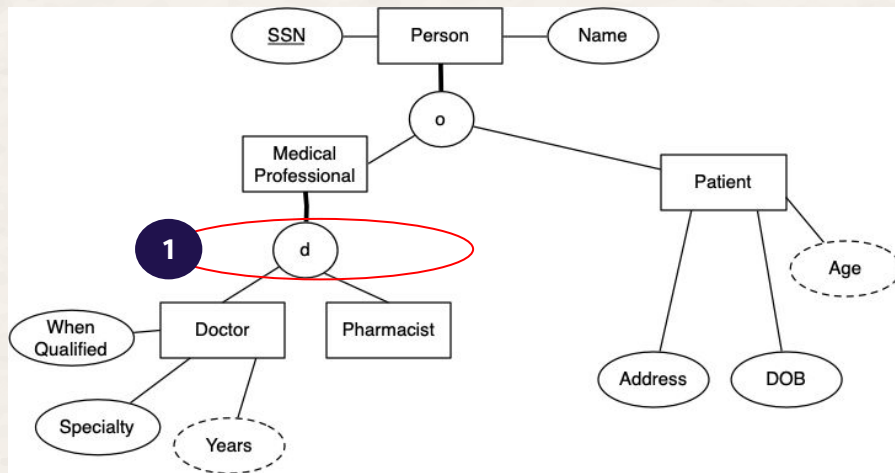
Relations correspond to **sets of ER entities**

Relations also correspond **ER relationships**



ER Diagram - Recap & Warm-up

Try to describe!



Relational Modelling – Recap & Warm-up

Try to answer!

What does a relation have?

- Relation name + a set of attributes

What does an attribute have?

- Attribute name + domain

What is a tuple?

- a list of values

$(2,3) = (3,2)$?

- $(2,3) \neq (3,2)$

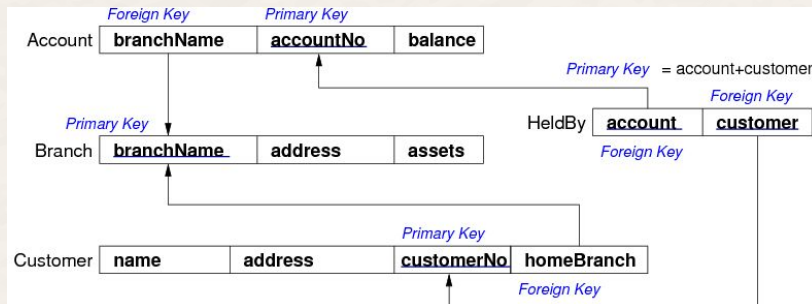
What is an instance of a relation?

- a set of tuples

$\{(a,b), (c,d)\} = \{(c,d), (a,b)\}$?

- $\{(a,b), (c,d)\} = \{(c,d), (a,b)\}$

How do we deal with referential integrity constraints?





ER-Relational Mapping

a useful strategy for database design:

- perform initial data modelling using ER (conceptual-level modelling)
- transform conceptual design into relational model (implementation-level modelling)

A formal mapping exists for ER model \rightarrow Relational model.

This maps "structures"; but additional info is needed, e.g. concrete **domains** for attributes and other **constraints**



ER-Relational Mapping



Correspondences between relational and ER data models:

$\text{attribute(ER)} \cong \text{attribute(Rel)}$, $\text{entity(ER)} \cong \text{tuple(Rel)}$

$\text{entity set(ER)} \cong \text{relation(Rel)}$, $\text{relationship(ER)} \cong \text{relation(Rel)}$

Differences between relational and ER models:

Rel uses **relations** to model entities and relationships

Rel has **no composite or multi-valued attributes** (only atomic)

Rel has **no object-oriented notions** (e.g. subclasses, inheritance)




Mapping *Strong Entities*



An entity set E with atomic attributes a_1, a_2, \dots, a_n

maps to

A relation R with attributes (columns) a_1, a_2, \dots, a_n



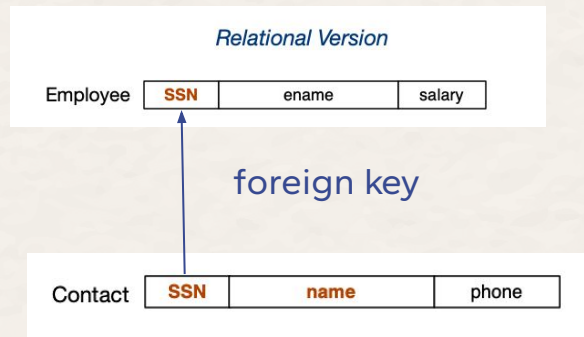

Note: the key is preserved in the mapping.

Mapping *Weak Entities*



What should the “Employee” relation look like?

What about “Contact”?



Mapping *N:M Relationships*



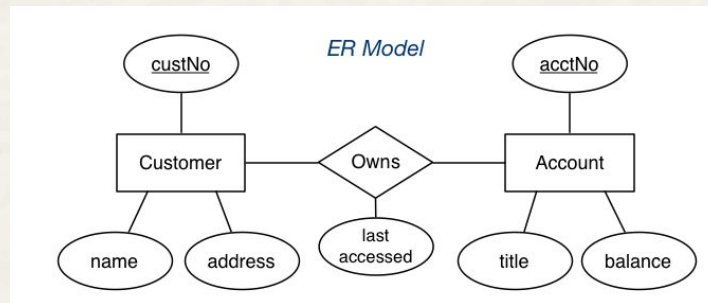
How to map N:M (many-to-many) relationships?

First, let's come up with the relations/tables for the entities.

Relational Version

Customer	custNo	name	address
----------	---------------	------	---------

Account	acctNo	title	balance
---------	---------------	-------	---------



Then, let's come up with the relation/table for the relationship.

Owns	acctNo	custNo	lastAccessed
------	---------------	---------------	--------------



primary keys + attributes

Mapping 1:N Relationships

How to map 1:N (1-to-many) relationships?

Is it similar to N:M relationships?

Customer


custNo	name	address
--------	------	---------

Branch

branchNo	address	assets
----------	---------	--------

HasHome


custNo	branchNo	joined
--------	----------	--------



Is HasHome really needed as a separate table? Can we merge it into another table? If so, which one should we use?

of HasHome depends on which entity?

of HasHome = # of Customer



Mapping 1:1 Relationships

How to map 1:1 (1-to-1) relationships?

Is it similar to N:M relationships?

Manager


empNo	name	salary
-------	------	--------

Branch

branchNo	address	assets
----------	---------	--------

Manages

empNo	branchNo
-------	----------



Can we merge the relationship table?
like what we did for 1:N relationship.


Relational Version

Manager	empNo	name	salary	branchNo
Branch	branchNo	address	assets	



Mapping 1:1 Relationships (cont)

If there is no reason to favour one side of the relationship ...



Relational Version #1

T	tid	a	sid	b
---	-----	---	-----	---

S	sid	c
---	-----	---

Relational Version #2

T	tid	a
---	-----	---

S	sid	c	tid	b
---	-----	---	-----	---





Mapping *n-way Relationships*

Relationship mappings above assume binary relationship.

If multiple entities are involved:

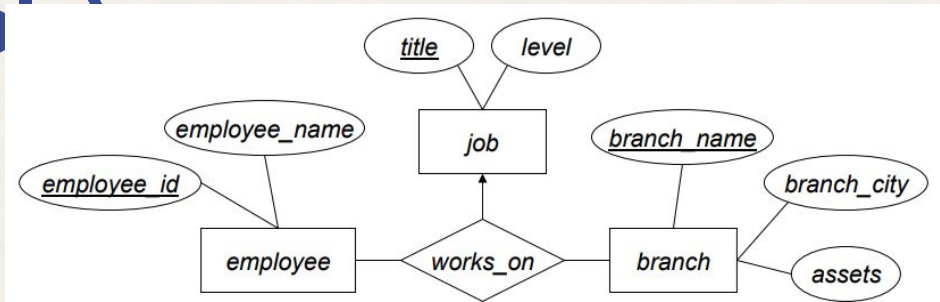
- $n:m$ generalises naturally to $n:m:p:q$
 - include foreign key for each participating entity
 - include any other attributes of the relationship
- other multiplicities (e.g. $1:n:m$) ...
 - need to be mapped the same as $n:m:p:q$
 - so not quite an accurate mapping of the ER

Some people advocate converting n -way relationships into:

- a new entity, and a set of n binary relationships



Mapping n -way Relationships (example)



What tables can we have for the entities?

```
job(title, level)
employee(employee_id, employee_name)
branch(branch_name, branch_city, assets)
```

What should the “works_on” table have as columns? (what attributes)

```
works_on(employee_id, branch_name, title)
```

Note: title is not included in the primary key here.



Mapping *n-way Relationships* (keys)

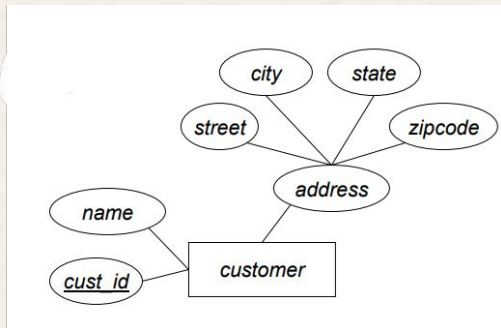


- If no-arrow (many-to-many):
 - primary key for the relationship is the union of all participating entity-sets' primary keys
- If one-arrow (one-to-many):
 - primary key for the relationship is the union of primary keys of entity-sets without an arrow
- Don't allow more than one arrow!
 - may think about redesigning the ER diagram



Mapping *Composite Attributes*

Composite attributes are mapped by concatenation or flattening.




customer(cust_id, name, street, city, state, zipcode)



Mapping *Multi-valued Attributes* (MVAs)

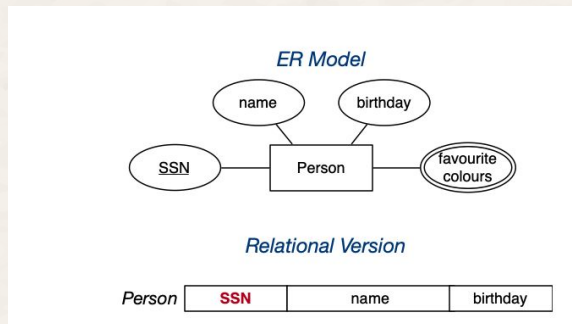
MVAs are mapped by a new table linking values to their entity.



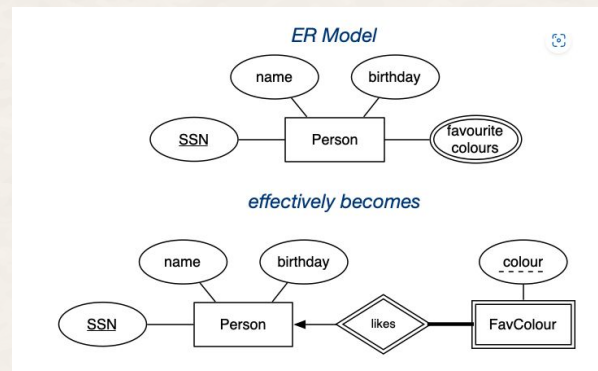
```
customer(cust_id, name)
cust_addrs(cust_id, street, city, state, zipcode)
```

Mapping *Multi-valued Attributes* (MVAs)

Another example.



What should the “favColors” table have?





Mapping *Subclasses*

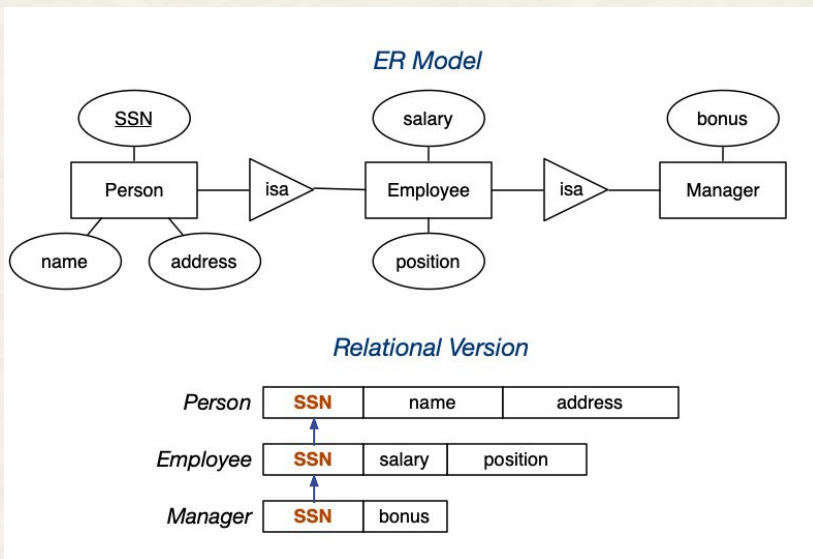
Three different approaches to mapping subclasses to tables:

- ER style
 - each entity becomes a separate table,
 - containing attributes of subclass + FK to superclass table
- object-oriented
 - each entity becomes a separate table,
 - inheriting all attributes from all superclasses
- single table with nulls
 - whole class hierarchy becomes one table,
 - containing all attributes of all subclasses (null, if unused)

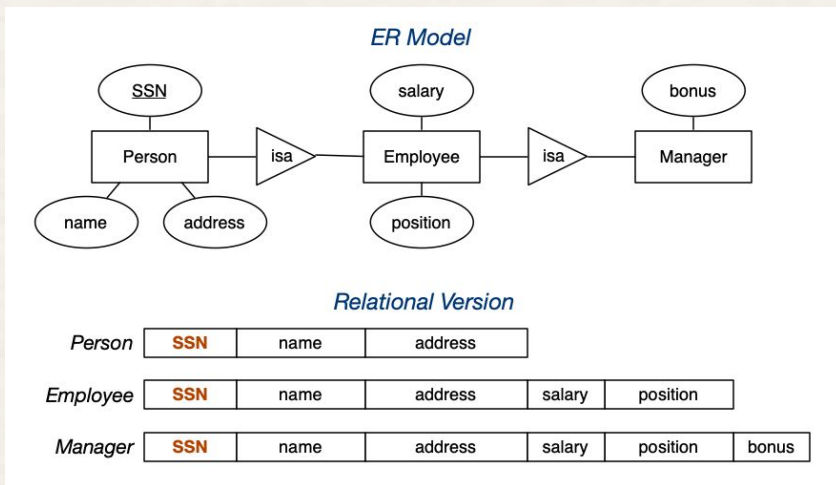
Which mapping is best depends on how data is to be used.



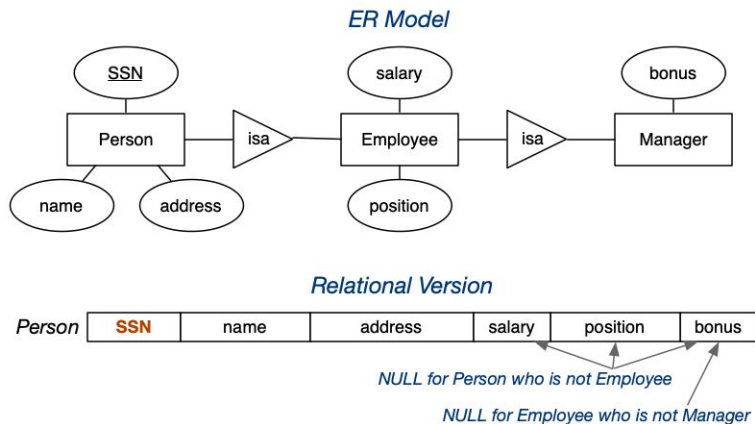
Mapping *Subclasses* - *ER style*



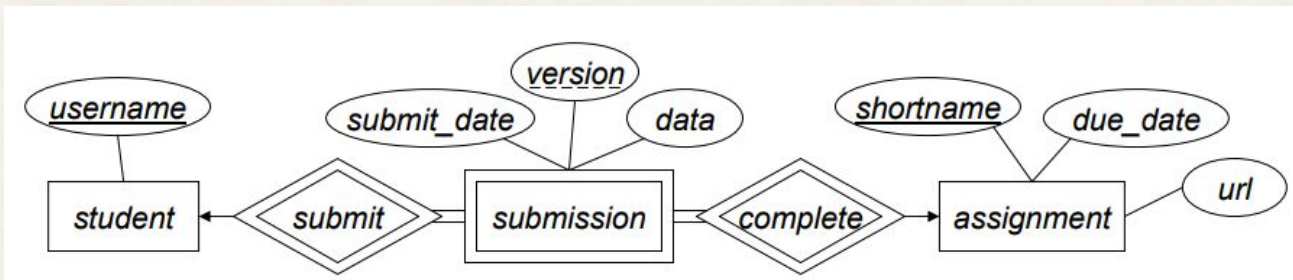
Mapping *Subclasses* - OO style



Mapping *Subclasses* - *single table style*



Exercise: Mapping Weak Entity



How to describe the two relations?

Every submission is submitted by **one** student.
A student **may** submit **one or more** submissions.

What are the tables for the strong entities?

```

student(username)
assignment(shortname, due_date, url)
  
```

What's the table for the weak entities?

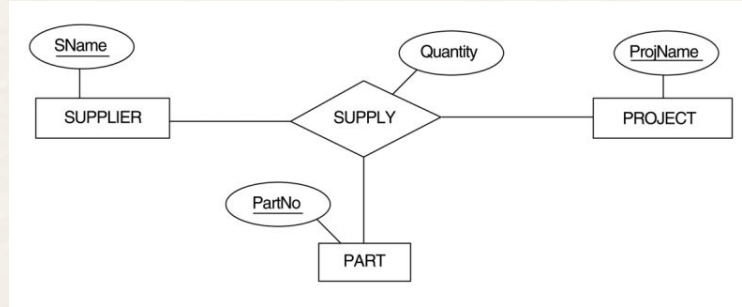
```

submission(username, shortname, version,
submit_date, data)
  
```

Every submission completes **one** assignment.
An assignment **may** be completed by **one or more** submissions.



Exercise: n-ary relationships



What are the tables for the entities?

SUPPLIER	
<u>SNAME</u>	...

PROJECT	
<u>PROJNAME</u>	...

PART	
<u>PARTNO</u>	...

What's the table for the relationship?

SUPPLY			
<u>SNAME</u>	<u>PROJNAME</u>	<u>PARTNO</u>	QUANTITY

Outline



- Announcements
- ER-Relational Mapping
- **SQL Introduction**





SQL vs Relational Model

- The relational model is a formal system for
 - describing data (relations, tuples, attributes, domains, constraints)
 - manipulating data (relational algebra ... covered elsewhere)
- SQL is a "programming" language for
 - describing data (tables, rows, fields, types, constraints)
 - manipulating data (query language)

More details:

https://runestone.academy/ns/books/published/practical_db/PART3_RELATIONAL_DATABASE_THEORY/04-sql-vs-theory/sql-vs-theory.html



SQL History



Developed at IBM in the mid-1970's (System-R)

- Standardised in 1986, and then in 1989, 1992, 1999, 2003, ... 2019
- Many database management systems (DBMSs) have been built around SQL
 - System-R, Oracle, Ingres, DB2, PostgreSQL, MySQL, SQL-server, SQLite, ...

DBMSs vs the standard

- all DBMSs implement a subset of the 1999 standard (aka SQL3)
- all DBMSs implement proprietary extensions to the standard

Conforming to standard should ensure portability of database applications



SQL Intro



SQL has several sub-languages ...

- **meta-data** definition language (e.g. create table, etc.)
- **meta-data** update language (e.g. alter table, drop table)
- **data** update language (e.g. insert, update, delete)
- **query** language (e.g. select ... from ... where, etc.)

Meta-data languages manage the **database schema**

Data update language manages **sets of tuples**



SQL Intro



Syntax-wise, SQL is similar to other programming languages

- has keywords, identifiers, constants, operators
- but strings are different to most PLs
 - '...' are constant strings, e.g. 'a', 'abc123', 'John's bag'
 - "..." allow non-alpha chars in identifiers and make id's case-sensitive

In the **standard**, **all non-quoted identifiers** map to **all upper-case**

e.g. BankBranches = bankbranches are treated as BANKBRANCHES

In **PostgreSQL**, **all non-quoted identifiers** map to **all lower-case**

e.g. BankBranches = BANKBRANCHES are treated as bankbranches



In all standards-adhering DBMSs, different quoted identifiers are different
"BankBranches" ≠ "bankbranches" ≠ "BANKBRANCHES"

SQL Syntax in a Nutshell

SQL definitions, queries and statements are composed of:

comments ... `--` comments to end of line

identifiers ... similar to regular programming languages

keywords ... a large set (e.g. `CREATE`, `DROP`, `TABLE`)

data types ... small set of basic types (e.g. integer, date)

operators ... similar to regular programming languages

constants ... similar to regular programming languages

- Similar means "often the same, but not always" ...

SQL programmers be like




`SELECT * FROM`

Select * From

select * from

SeLEct * fRoM



Try it: [PostgreSQL - OneCompiler - Write, run and share PostgreSQL code online](#)



SQL Syntax in a Nutshell

Comments: everything after `--` is a comment

Identifiers: alphanumeric (`a la C`), but also "An Identifier"

Reserved words: many e.g. `CREATE`, `SELECT`, `TABLE`, ...

Reserved words cannot be used identifiers unless quoted e.g. "table"

Strings: e.g. `'a string'`, `'don't ask'`, but no `'\n'` (use `e'\n'`)

Numbers: like C, e.g. `1`, `-5`, `3.14159`, ...

Try it: [PostgreSQL - OneCompiler - Write, run and share PostgreSQL code online](#)



Types: `integer`, `float`, `char(n)`, `varchar(n)`, `date`, ...

Operators: `=`, `<>`, `<`, `<=`, `>`, `>=`, `AND`, `OR`, `NOT`, ...

Names in SQL



Identifiers denote:

- database objects such as tables, attributes, views, ...
- meta-objects such as types, functions, constraints, ...

Naming conventions that I (try to) use in this course:

- **relation names:** e.g. Branches, Students, ... (use plurals)
- **attribute names:** e.g. name, code, firstName, ...
- **foreign keys:** named after either or both of
 - table being referenced e.g. staff or staff_id, ...
 - relationship being modelled e.g. teaches, ...

We initially write SQL keywords in all upper-case in slides.



Types/Constants in SQL - Numeric



Name	Storage Size	Description	Range
smallint	2 bytes	small-range integer	-32768 to +32767
integer	4 bytes	typical choice for integer	-2147483648 to +2147483647
bigint	8 bytes	large-range integer	-9223372036854775808 to +9223372036854775807
decimal	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
numeric	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
real	4 bytes	variable-precision, inexact	6 decimal digits precision
double precision	8 bytes	variable-precision, inexact	15 decimal digits precision
smallserial	2 bytes	small autoincrementing integer	1 to 32767
serial	4 bytes	autoincrementing integer	1 to 2147483647
bigserial	8 bytes	large autoincrementing integer	1 to 9223372036854775807

arbitrary numeric data:

```
NUMERIC(precision, scale)
```

e.g., the number 23.5141 has a precision of 6 and a scale of 4. Integers can be considered to have a scale of zero.

<https://www.postgresql.org/docs/current/datatype-numeric.html>





Types/Constants in SQL - Char/String

Name	Description
<code>character varying(n)</code> , <code>varchar(n)</code>	variable-length with limit
<code>character(n)</code> , <code>char(n)</code> , <code>bpchar(n)</code>	fixed-length, blank-padded
<code>bpchar</code>	variable unlimited length, blank-trimmed
<code>text</code>	variable unlimited length

PostgreSQL provides extended strings containing \ escapes, e.g:

`E'\n'` `E'O\Brien'` `E'[A-Z]{4}\\d{4}'` `E'John'`

Try it: [PostgreSQL - OneCompiler - Write, run and share PostgreSQL code online](#)

Type-casting via `Expr::Type` (e.g. `'10'::integer`)



<https://www.postgresql.org/docs/current/datatype-character.html>



Types/Constants in SQL - Logical

Logical type: `BOOLEAN`, `TRUE` and `FALSE` (or `true` and `false`)

PostgreSQL also allows `'t'`, `'true'`, `'yes'`, `'f'`, `'false'`, `'no'`

```
-- create
CREATE TABLE EMPLOYEE (
  empId INTEGER PRIMARY KEY,
  name TEXT NOT NULL,
  dept TEXT NOT NULL,
  retired BOOLEAN default 'no'
);

-- insert
INSERT INTO EMPLOYEE VALUES (0001, 'Clark', 'Sales');
INSERT INTO EMPLOYEE VALUES (0002, 'Dave', 'Accounting');
INSERT INTO EMPLOYEE VALUES (0003, 'Ava', 'Sales', true);

-- fetch
SELECT * FROM EMPLOYEE WHERE retired = false;
```

Try it: [PostgreSQL - OneCompiler - Write, run and share PostgreSQL code online](#)





Types/Constants in SQL - Time

Name	Storage Size	Description	Low Value	High Value	Resolution
timestamp [(p)] [without time zone]	8 bytes	both date and time (no time zone)	4713 BC	294276 AD	1 microsecond
timestamp [(p)] with time zone	8 bytes	both date and time, with time zone	4713 BC	294276 AD	1 microsecond
date	4 bytes	date (no time of day)	4713 BC	5874897 AD	1 day
time [(p)] [without time zone]	8 bytes	time of day (no date)	00:00:00	24:00:00	1 microsecond
time [(p)] with time zone	12 bytes	time of day (no date), with time zone	00:00:00+1559	24:00:00-1559	1 microsecond
interval [fields] [(p)]	16 bytes	time interval	-178000000 years	178000000 years	1 microsecond

Subtraction of timestamps yields an interval



<https://www.postgresql.org/docs/current/datatype-datetime.html>

Types/Constants in SQL – Others



PostgreSQL also has a range of non-standard types, e.g.

- geometric (point/line/...), currency, IP addresses, JSON, XML, objectIDs, ...
- non-standard types typically use string literals ('...') which need to be interpreted

<https://www.postgresql.org/docs/current/datatype.html>



Types/Constants in SQL - User-defined



-- domains: constrained version of existing type

```
CREATE DOMAIN Name AS Type CHECK ( Constraint )
```

-- tuple types: defined for each table

```
CREATE TYPE Name AS ( AttrName AttrType, ... )
```

-- enumerated type: specify elements and ordering

```
CREATE TYPE Name AS ENUM ( 'Label', ... )
```



Types/Constants in SQL - User-defined (Examples)



-- positive integers

```
CREATE DOMAIN PosInt AS integer CHECK (value > 0);
```

-- a UNSW course code

```
CREATE DOMAIN CourseCode AS char(8)  
CHECK (value ~ '[A-Z]{4}[0-9]{4}');
```

-- a UNSW student/staff ID

```
CREATE DOMAIN ZID AS integer  
CHECK (value between 1000000 and 9999999);
```

-- standard UNSW grades (FL,PS,CR,DN,HD)

```
CREATE DOMAIN Grade AS char(2)  
CHECK (value in ('FL','PS','CR','DN','HD'));
```

-- or

```
CREATE TYPE Grade AS ENUM ('FL','PS','CR','DN','HD');
```



Tuple and Set Literals



Tuple and set constants are both written as:

```
( val1, val2, val3, ... )
```

The correct interpretation is worked out from the context.

```
INSERT INTO Student(studID, name, degree)
```

```
VALUES (2177364, 'Jack Smith', 'BSc')
```

-- tuple literal

```
CONSTRAINT CHECK gender IN ('male','female','unspecified')
```

-- set literal



Thank you!