

COMP3311 24T1

Database Systems

week 2 - 2



Outline



- **Announcements**
- SQL Expressions
- SQL Meta-Data Definition Language
- Managing RDBMS
- Mapping ER->Rel->SQL



Announcement 1



- Quiz 1

- [Activities | COMP3311 24T1 | WebCMS3 \(unsw.edu.au\)](#)

- You can try as many times as possible

- before

midnight Saturday, 24th Feb (11:59:59 pm),
TRY to do it before 8pm Friday



Power Outage (CSE building): 8pm Friday - 8pm
Saturday



Announcement 2

How to learn during this stage of the course? 🤔

A long time ago, in a university far far away ...

When I was an undergraduate student ...

- Don't be afraid of the overwhelming concepts/syntax/keywords to remember
- The lectures are to “store” some **keywords** into your mind
- Practice A LOT! (lecture exercises, tutorials, assignment ...)
- Google A LOT! (with the **keywords** in your mind | stackoverflow, chatgpt, ed forum ...)



Outline



- Announcements
- **SQL Expressions**
- SQL Data Definition Language
- Managing RDBMS
- Mapping ER->Rel->SQL



Expressions in SQL



Expressions in SQL involve: **objects**, **constants**, **operators**

- objects are typically names of attributes (or PLpgSQL variables)
- operators may be symbols (e.g. `+`, `=`) or keywords (e.g. `between`)

SQL constants are similar to typical programming language constants

- integers: `123`, `-5`; floats: `3.14`, `1.0e-3`; boolean: `true`, `false`

But strings are substantially different

- `'...'` rather than `"..."`, no `\n`-like "escape" chars
- escape mechanisms: `'O'Brien'` or `E'O'Brien'` (non-standard)
- dollar quoting: `$$O'Brien$$` or `tagO'Brientag`



SQL Operators



Comparison operators are defined on all types:

- `<` `>` `<=` `>=` `=` `<>`

In PostgreSQL, `!=` is a synonym for `<>` (but there's no `==`)

Boolean operators `AND`, `OR`, `NOT` are also available

- Note `AND`, `OR` are not "short-circuit" in the same way as C's `&&`, `||`

Most data types also have type-specific operations available

- String comparison (e.g. `str1 < str2`) uses dictionary order



<https://www.postgresql.org/docs/16/functions.html>



SQL Operators (string-matching)

SQL provides pattern matching for strings via `LIKE` and `NOT LIKE`

- `%` matches anything (cf. regexp `.*`)
- `_` matches any single char (cf. regexp `.`)

Examples:

- `name LIKE 'Ja%'` name begins with 'Ja'
- `name LIKE '_i%'` name has 'i' as 2nd letter
- `name LIKE '%o%o%'` name contains two 'o's
- `name LIKE '%ith'` name ends with 'ith'
- `name LIKE 'John'` name equals 'John'

PostgreSQL also supports case-insensitive matching: `ILIKE`



SQL Operators (string-matching)



PostgreSQL provides regexp-based pattern matching via `~` and `!~`

Examples (using POSIX regular expressions):

- `name ~ '^Ja'` name begins with 'Ja'
- `name ~ '^i'` name has 'i' as 2nd letter
- `name ~ '.*o.*o.*'` name contains two 'o's
- `name ~ 'ith$'` name ends with 'ith'
- `name ~ 'John'` name contains 'John'

Also provides case-insensitive matching via `~*` and `!~*`



SQL Operators (string)



Other operators/functions for string manipulation:

`str1 || str2` ... return concatenation of str1 and str2

`lower(str)` ... return lower-case version of str

`substring(str,start,count)` ... extract substring from str

Etc. ...

Note that above operations are null-preserving (strict):

if any operand is `NULL`, result is `NULL`



SQL Operators (arith)



Arithmetic operations:

`+` `-` `*` `/` `abs` `ceil` `floor` `power` `sqrt` `sin` etc.

Aggregations "summarize" a column of numbers in a relation:

- `count(attr)` ... number of rows in attr column
- `sum(attr)` ... sum of values for attr
- `avg(attr)` ... mean of values for attr
- `min/max(attr)` ... min/max of values for attr

Note: **count** applies to columns of non-numbers as well.



The NULL value

Expressions containing NULL generally yield NULL.

However, boolean expressions use three-valued logic:



a	b	a AND b	a OR b
T	T	T	T
T	F	F	T
T	NULL	NULL	T
F	T	F	T
F	F	F	F
F	NULL	F	NULL
NULL	NULL	NULL	NULL



The NULL value (cont)



Important consequence of NULL behaviour ...

These expressions do not work as (might be) expected:

`x = NULL` `x <> NULL`

Both return NULL regardless of the value of `x`

Can only test for NULL using:



`x IS NULL` `x IS NOT NULL`





Conditional Expressions

Other ways that SQL provides for dealing with **NULL**:

`coalesce(val1,val2,...valn)`

- returns first non-null value `vali` (or **NULL** if all values are **NULL**).

E.g. `select coalesce(mark, ...) from Marks ...`

`nullif(val1,val2)`

- takes two values and returns the first value, except it returns **NULL** if `val1` is equal to `val2`

E.g. `nullif(mark,'??')`





Conditional Expressions

SQL also provides a generalised conditional expression:

CASE

WHEN test₁ **THEN** result₁

WHEN test₂ **THEN** result₂

...

ELSE result_n

END

E.g. `case when mark >= 85 then 'HD' ... else '??' end`

Tests that yield **NULL** are treated as **FALSE**

If no **ELSE**, and all tests fail, **CASE** yields **NULL**



Outline



- Announcements
- SQL Expressions
- **SQL Meta-Data Definition Language**
- Managing RDBMS
- Mapping ER->Rel->SQL



SQL Intro – Recap



SQL has several sub-languages ...

- **meta-data** definition language (e.g. create table, etc.)
- **meta-data** update language (e.g. alter table, drop table)
- **data** update language (e.g. insert, update, delete)
- **query** language (e.g. select ... from ... where, etc.)

Meta-data languages manage the **database schema**

Data update language manages **sets of tuples**



Relational Data Definition



In order to give a relational data model, we need to:

- describe **tables**
- describe **attributes** that comprise tables
- describe any **constraints** on the data

A **relation schema** defines an individual table

- table name, attribute names, attribute domains, keys, etc.

A **database schema** is a collection of relation schemas that

- defines the structure the whole database
- additional constraints on the whole database





Defining a Relation Schema

Tables (relations) are described using:

```
CREATE TABLE TableName (  
    attribute1 domain1 constraints1,  
    attribute2 domain2 constraints2,  
    ...  
    table-level constraints, ...  
)
```

This SQL statement ..

- defines the table schema (adds it to database meta-data)
- creates an empty instance of the table (zero tuples)



Tables are removed via `DROP TABLE TableName;`

Defining a Relation Schema



Example: defining the Students table ...

```
CREATE TABLE Students (  
  zid serial,  
  family varchar(40),  
  given varchar(40) NOT NULL,  
  d_o_b date NOT NULL,  
  gender char(1) CHECK (gender in ('M','F','N')),  
  degree integer,  
  PRIMARY KEY (zid),  
  FOREIGN KEY (degree) REFERENCES Degrees(did)  
);
```

```
CREATE TABLE Students (  
  zid serial PRIMARY KEY, -- only if it's a single attribute key  
  family varchar(40),  
  given varchar(40) NOT NULL,  
  d_o_b date NOT NULL,  
  gender char(1) CHECK (gender in ('M','F','N')),  
  degree integer REFERENCES Degrees(did)  
);
```



A primary key attribute is implicitly defined to be **UNIQUE** and **NOT NULL**



Data Integrity



Defining tables as above affects behaviour of DBMS when changing data

Constraints and types ensure that integrity of data is preserved

- no duplicate keys
- no "dangling references"
- all attributes have valid values
- etc. etc. etc.

Preserving data integrity is a **critical** function of a DBMS.



Defining Keys



Primary keys:

- if PK is one attribute, can define as attribute constraint
- if PK is multiple attributes, must define in table constraints
- PK implies **NOT NULL, UNIQUE** for all attributes in key

Foreign keys:

- if FK is one attribute, can define as attribute constraint
- can omit **FOREIGN KEY** keywords in attribute constraint
- if FK has multiple attributes, must define as a single table constraint
- should always specify corresponding PK attribute in FK constraint



Defining Keys (cont)



Defining primary keys assures **entity integrity**

- must give values for all attributes in the primary key

For example this insertion would fail ...

```
INSERT INTO Enrolments(student,course,mark,grade)  
VALUES (5123456, NULL, NULL, NULL);
```

because no course was specified; but mark and grade can be NULL

Defining primary keys assures **uniqueness**

- cannot insert a tuple which contains an existing PK value



Defining Keys (cont)



Defining foreign keys assures **referential integrity**.

On insertion, cannot add a tuple where FK value does not exist as a PK

For example, this insert would fail ...

```
INSERT INTO Accounts(acctNo, owner, branch, balance)  
VALUES ('A-123', 765432, 'Kensington', 5000);
```

if there is no customer with id 765432 or no branch Kensington





Attribute Value Constraints

NOT NULL and **UNIQUE** are special constraints on attributes.

SQL has a general mechanism for specifying attribute constraints

attrName type CHECK (Condition)

Condition is a boolean expression and can involve other attributes, relations and sub-queries.

```
CREATE TABLE Example
(
  gender char(1) CHECK (gender IN ('M','F')),
  Xvalue integer NOT NULL,
  Yvalue integer CHECK (Yvalue > Xvalue),
  Zvalue float CHECK (Zvalue >
    (SELECT MAX(price)
     FROM Sells)
  )
);
```

(but many RDBMSs (e.g. Oracle and **PostgreSQL**) don't allow subqueries in CHECK)





Named Constraints

A constraint in an SQL table definition can (optionally) be named via

```
CONSTRAINT constraintName constraint
```

Example:

```
CREATE TABLE Example  
(  
  gender char(1) CONSTRAINT GenderCheck  
    CHECK (gender IN ('M','F','N')),  
  Xvalue integer NOT NULL,  
  Yvalue integer CONSTRAINT XYOrder  
    CHECK (Yvalue > Xvalue)  
);
```



Try it: [PostgreSQL - OneCompiler - Write, run and share PostgreSQL code online](#)

reason:

<https://stackoverflow.com/questions/1397440/what-is-the-purpose-of-constraint-naming#:~:text=By%20naming%20the%20constraints%20you%20can%20differentiate%20violations,react%20differently%20depending%20on%20which%20constraint%20was%20violated.>

Outline



- Announcements
- SQL Expressions
- SQL Meta-Data Definition Language
- **Managing RDBMS**
- Mapping ER->Rel->SQL




What is an RDBMS?



- A relational database management system (RDBMS) is
 - software designed to support large-scale data-intensive applications
 - allowing high-level description of data (tables, constraints)
 - with high-level access to the data (relational model, SQL)
 - providing efficient storage and retrieval (disk/memory management)
 - supporting multiple simultaneous users (privilege, protection)
 - doing multiple simultaneous operations (transactions, concurrency)
 - maintaining reliable access to the stored data (backup, recovery)
- Note: databases provide persistent storage of information (even for redis !)



Postgres Structure



Managing Databases



Shell commands to create/remove databases:

- `createdb dbname` ... create a new totally empty database
- `dropdb dbname` ... remove all data associated with a DB

Shell commands to dump/restore database contents:

```
pg_dump dbname > dumpfile
```

```
psql dbname -f dumpfile
```

(Database dbname is typically created just before restore)





Managing Tables

SQL statements:

- **CREATE** TABLE table (Attributes+Constraints)
- **ALTER** TABLE table TableSchemaChanges
- **DROP** TABLE table(s) [CASCADE]
- **TRUNCATE** TABLE table(s) [CASCADE]

(All conform to SQL standard, but all also have extensions)

DROP.CASCADE also drops objects which depend on the table

- objects could be tuples or views, but not whole tables

TRUNCATE.CASCADE truncates tables which refer to the table



<https://www.postgresql.org/docs/current/ddl.html>



Managing Tuples

SQL statements:

- **INSERT** INTO table (Attrs) VALUES Tuple(s)
- **DELETE** FROM table WHERE condition
- **UPDATE** table SET AttrValueChanges WHERE condition

Attrs = (attr₁, ... attr_n) Tuple = (val₁, ... val_n)

AttrValueChanges is a comma-separated list of:

- attrname = expression
- Each list element assigns a new value to a given attribute.



<https://www.postgresql.org/docs/current/dml.html>

Outline



- Announcements
- SQL Expressions
- SQL Meta-Data Definition Language
- Managing RDBMS
- **Mapping ER->Rel->SQL**




Mapping *Strong Entities*



An entity set E with atomic attributes a_1, a_2, \dots, a_n

maps to

A relation R with attributes (columns) a_1, a_2, \dots, a_n



SQL Version


```
create table Students (  
  stuID integer primary key,  
  name text not null,  
  degree char(4)  
);
```



Note: the key is preserved in the mapping.


Mapping *Weak Entities*

- **foreign keys:** named after either or both of
 - **table being referenced** e.g. staff or staff_id, ...
 - **relationship being modelled** e.g. teaches, ...



```
create table Employees (
  SSN text primary key,
  ename text,
  salary currency
);
```

```
create table Contacts (
  relatedTo text not null, -- total participation
  name text, -- not null implied by PK
  phone text not null,
  primary key (relatedTo, name),
  foreign key (relatedTo) references Employees (ssn)
);
```



foreign key




Mapping *N:M Relationships*

Relational Version

Customer	custNo	name	address
Account	acctNo	title	balance

Owns	acctNo	custNo	lastAccessed
------	---------------	---------------	--------------



ER Model



```

create table Customers (
    custNo serial primary key,
    name text not null,
    address text -- don't need to know customer's address
);
create table Accounts (
    acctNo char(5) check (acctNo ~ '[A-Z]-[0-9]{3}'),
    title text not null, -- acctNos are like 'A-123'
    balance float default 0.0,
    primary key (acctNo)
);
create table Owns (
    customer_id integer references Customers(custNo),
    account_id char(5) references Accounts(acctNo),
    last_accessed timestamp,
    primary key (customer_id, account_id)
);
    
```


Mapping 1:N Relationships



```
create table Branches (
    branchNo serial primary key,
    address text not null,
    assets currency
);
create table Customers (
    custNo serial primary key,
    name text not null,
    address text,
    hasHome integer not null, -- total participation
    joined date not null,
    foreign key (hasHome) references Branches(branchNo)
);
```

Mapping 1:1 Relationships

Relational Version


Manager	empNo	name	salary	branchNo
Branch	branchNo	address	assets	

```

create table Branches (
  branchNo serial primary key,
  address  text not null,
  assets   currency
);
-- a new branch
-- may have no accounts

create table Managers (
  empNo serial primary key,
  name   text not null,
  salary currency not null, -- when first employed,
                                -- must have a salary
  manages integer not null, -- total participation
  foreign key (manages) references Branches(branchNo)
);
  
```

ER Model




What if there's a total participation for "Branch"?

If both entities have total participation, cannot express this in SQL except by putting a (redundant) **not null** foreign key in one table.

i.e. put a foreign key of **empNo** in the branch table.

Mapping *n*-way Relationships



What if a customer can only have one account?
(i.e., there's an arrow pointing to account)
What's the primary key?




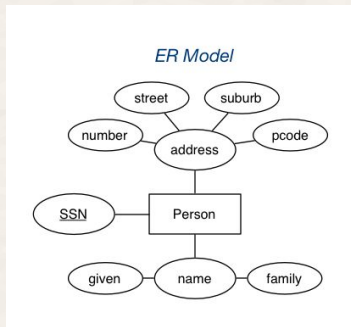
customer_id, atm_id

```

create table Customers (
    custNo serial primary key, ...
);
create table Accounts (
    acctNo char(5) ... primary key, ...
);
create table ATMs (
    atmID serial primary key,
    cash currency check (cash >= 0),
    location text not null
);
create table Withdrawal (
    customer_id integer references Customers(custNo),
    account_id char(5) references Accounts(acctNo),
    atm_id integer references ATMs(atmID),
    amount currency not null,
    when timestamp default now(),
    primary key (customer_id, atm_id)
);
  
```

Mapping Composite Attributes

Composite attributes are mapped by concatenation or flattening.




```
create table People (
  ssn      integer primary key,
  given    text not null,
  family   text,
  number   integer not null,
  street   text not null,
  suburb   text not null,
  pcode    char(4) not null check (pcode ~ '[0-9]{4}')
);
```



Mapping *Multi-valued Attributes* (MVAs)

MVAs are mapped by a new table linking values to their entity.



```
create table People (  
    ssn      integer primary key,  
    name     text not null,  
    birthday date  
);  
create table FavColour (  
    person_id integer references People(ssn),  
    colour    text,  
    primary key (person_id, colour)  
);
```





Mapping *Subclasses*

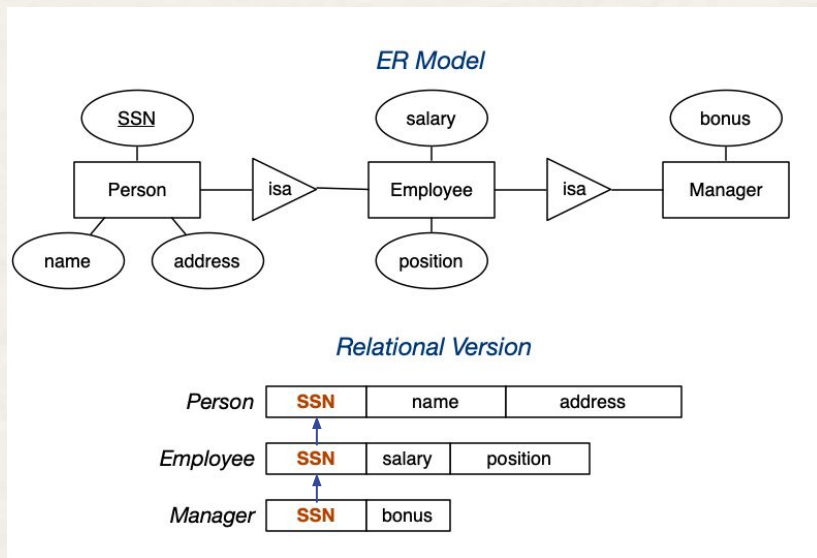
Three different approaches to mapping subclasses to tables:

- ER style
 - each entity becomes a separate table,
 - containing attributes of subclass + FK to superclass table
- object-oriented
 - each entity becomes a separate table,
 - inheriting all attributes from all superclasses
- single table with nulls
 - whole class hierarchy becomes one table,
 - containing all attributes of all subclasses (null, if unused)

Which mapping is best depends on how data is to be used.




Mapping Subclasses - ER style



```
create table People (  
    ssn    integer primary key,  
    name   text not null,  
    address text  
);  
create table Employees (  
    person_id integer primary key,  
    salary    currency not null,  
    position  text not null,  
    foreign key (person_id) references People(ssn)  
);  
create table Managers (  
    employee_id integer primary key,  
    bonus       currency,  
    foreign key (employee_id)  
                references Employees(person_id)  
);
```

Mapping Subclasses - OO style

ER Model



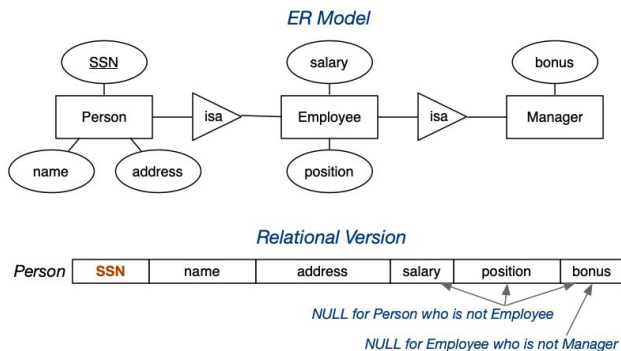
Relational Version

Person	SSN	name	address			
Employee	SSN	name	address	salary	position	
Manager	SSN	name	address	salary	position	bonus

```

create table People (
  ssn    integer primary key,
  name   text not null,
  address text
);
create table Employees (
  ssn    integer primary key,
  name   text not null,
  address text
  salary currency not null,
  position text not null
);
create table Managers (
  ssn    integer primary key,
  name   text not null,
  address text
  salary currency not null,
  position text not null,
  bonus  currency
);
  
```


Mapping Subclasses - single table style



```
create table People (
  ssn      integer primary key,
  ptype    char(1) not null
           check (ptype in ('P','E','M')),
  name     text not null,
  address  text,
  salary   currency,
  position text,
  bonus    currency,
  constraint subclasses check
    ((ptype = 'P' and salary is null
     and position is null and bonus is null)
     or
     (ptype = 'E' and salary is not null
     and position is not null and bonus is null)
     or
     (ptype = 'M' and salary is not null
     and position is not null and bonus is not null))
);
```

Thank you!